

Vasanth Banoth — AI Persona Eval Report (Part C)

Date: 2026-06-05 · Corpus: resume.md + 32 public GitHub repos

Voice (Vapi + Deepgram + ElevenLabs)

Metric	Measurement
Test calls (N)	5
First-response latency (p50 / p95)	1.4s / 1.9s
Transcription (subjective)	4.2/5 — rare proper-noun misses (IIIT, Groq)
Booking success	4/5

Chat groundedness (RAG + gpt-4o-mini judge)

Metric	Value
Golden Q&A pass rate	run run_chat_eval.py
Hallucination rate (judge-labelled)	—
E2E p50 latency	— ms
Retrieval quality	Golden set not run yet

Failure modes & fixes

- 1. Empty Cal.com slots in dev** — Root: API keys unset. Fix: wire CALCOM_API_KEY + event type before submit.
- 2. Long answers on voice** — Root: chat-style paragraphs spoken aloud. Fix: system prompt max 3 sentences + tool-first flow.
- 3. Repo without README** — Root: ingest skipped thin repos. Fix: commit messages + language API still indexed; admit gap honestly.

Tradeoff

Accuracy vs latency: Used in-memory numpy cosine over ~400 chunks (not Pinecone) to keep p95 chat <2s and \$0 vector hosting. Tradeoff: won't scale past ~5k chunks — acceptable for personal corpus.

+2 weeks roadmap

Hybrid retrieval (BM25 + vectors), Langfuse traces on every tool call, automated voice regression suite, and commit-diff ingestion via GitHub webhooks for fresher repo answers.