

Vasanth Banoth — AI Persona Eval Report (Part C)

Date: 2026-06-07 · Corpus: resume.md + public GitHub repos

Voice (Vapi + Deepgram + ElevenLabs)

Metric	Measurement
Test calls (N)	5
First-response latency (p50 / p95)	1.4s / 1.9s
Transcription (subjective)	4.2/5 — rare proper-noun misses (IIIT, Groq)
Booking success	4/5

Chat groundedness (hybrid RAG + LLM judge)

Metric	Value
Golden Q&A pass rate	100%
Hallucination rate (judge-labelled)	0%
E2E p50 latency	2671 ms
Retrieval quality	7-case golden set; 7 distinct source tags hit

Failure modes & fixes

- 1. Cal.com unconfigured in dev** — Root: missing API keys. Fix: set `CALCOM_API_KEY` + `CALCOM_EVENT_TYPE_ID` on Render.
- 2. LLM quota exhausted** — Root: OpenAI billing. Fix: Groq fallback + corpus-only synthesize fallback (no invented facts).
- 3. Stale Next.js .next cache** — Root: build + dev conflict. Fix: `start_all.sh` clears `.next` before dev server.

Tradeoff

Accuracy vs latency: In-memory hybrid BM25 + BGE vectors over personal corpus (~35–400 chunks) instead of hosted Pinecone — keeps retrieval <200ms and \$0 vector hosting. Tradeoff: won't scale past ~5k chunks; fine for resume + GitHub scope.

+2 weeks roadmap

Langfuse traces on every tool call, automated voice regression from `TEST_SCRIPTS.md`, GitHub webhook re-ingest on push, and multilingual voice (Hindi code-switch) using Josh-AI-TASK learnings.